

# DataSci 510

## lesson 8

### unsupervised learning



## today's agenda

- unsupervised learning
- supervised vs unsupervised
- k-means clustering
- k-means vs. k-nearest neighbor
- k-means assumptions
- where k-means fails

# unsupervised learning

- the goal is to **find structure** in the data
- look at **unlabeled data** and find general patterns
- more subjective and difficult to evaluate and interpret, and hence it is far **less common** than supervised learning
- **clustering** is the most common example
  - k-means clustering
  - variable clustering / dimensionality reduction
  - word clouds (kind of)

## supervised vs unsupervised

- let's take **anomaly detection** as an example, such as **fraud detection**, **intrusion detection**, **health monitoring**
  1. if we have historical data where past anomalies are **labeled**, we can use **supervised learning** to predict future anomalies
  2. if the data is **unlabeled**, we can use **unsupervised learning** such as **k-means** to detect clusters that look suspicious
- **no free lunch**: labeling data can be expensive and time-consuming, but so is examining and interpreting clusters

## k-means characteristics


- there are **no labels**: k-means is **unsupervised**
- clusters are a **construct** we create, not something set in stone
- clusters can be hard to interpret
  - lots of **gray areas** when comparing clusters
  - k-means provides **hard clusters** but **soft clusters** are better
- there is a **supervised learning** algorithm that is very similar to k-means in how it works, called **k-nearest neighbors**
  - unlike k-means, it can be easily **evaluated**

## k-means clustering

- here we chose  $k = 10$
- we have two **numeric** features
- the white crosses are **cluster centroids**
- the colors show cluster **assignments**

image source: [[scikit-learn.org](https://scikit-learn.org)]


K-means clustering on the digits dataset (PCA-reduced data)  
Centroids are marked with white cross



## k-nearest neighbor

- $k$  is the number of neighbors to consider
- colors show the labels
- the colors of regions show decision boundaries
- larger  $k$  makes decision boundary smoother

image source: [[scikit-learn.org](https://scikit-learn.org)]



**notebook time**

**we return to the lecture later**



## k-means algorithm implementation

1. start with  $k$  random **centroids** in the **feature space**, preferably spread out well
2. calculate the **Euclidean distance** of every row to each of the  $k$  centroids
3. **assign** each row to whichever centroid it is closest to
4. **recalculate** cluster centroids
5. **repeat** steps 2 through 4 until results stabilize

## k-means assumption

- **Euclidean distance** means that
  - categorical data must be represented numerically
  - numeric data must be **normalized**
- we want to maximize variability **between clusters**
  - i.e. cluster centroids should be far away from each other
- we want to minimize variability **within clusters**
  - i.e. points belonging to the same cluster should be close to the centroid of their cluster

## numeric distance metrics

Let  $u = (u_1, u_2, \dots, u_n)$  and  $v = (v_1, v_2, \dots, v_n)$ .

Euclidean and Manhattan distance are part of a larger family called Minkowski distance:

- the **Euclidean distance**  $\text{dist}(u, v) = \sum_{i=1}^n (u_i - v_i)^2$  is also called the **L2-metric**
- the **Manhattan distance**  $\text{dist}(u, v) = \sum_{i=1}^n |u_i - v_i|$  is also called the **L1-metric**

Cosine similarity is a measure of the angle made between the vectors  $u$  and  $v$ . It is a good choice when directionality matters more than position (e.g. word vectors):

- the **cosine similarity** is given by  $\text{dist}(u, v) = \frac{u \cdot v}{\|u\| \cdot \|v\|}$  where the product at the numerator is a **dot product** but the product in the denominator is a simple product

## categorical distance metrics

- the **Hamming distance** lines up strings (same length) and counts the number of positions that don't match
- the **Levenshtein distance** (also called the **edit distance**) between any two strings is the total cost of the number of insertions (costs 1), deletions (costs 1), substitutions (costs 2) needed to convert one string into the other
- the **Jaccard index** measures the size of the **intersection** of characters divided by size of the **union** of characters:  $J(A, B) = 1 - \frac{|A \cap B|}{|A \cup B|}$ , for example  
 $J(\text{beer}, \text{bear}) = 1 - \frac{3}{4}$

## discussion

The **weighted Jaccard index** is based on the minimum number of times (denoted  $m_i$ ) that a letter appears in either word and the number of times it appears in both words **combined** (denoted  $M_i$ ). It is given by:

$$J'(A, B) = 1 - \frac{\sum m_i}{\sum M_i}$$

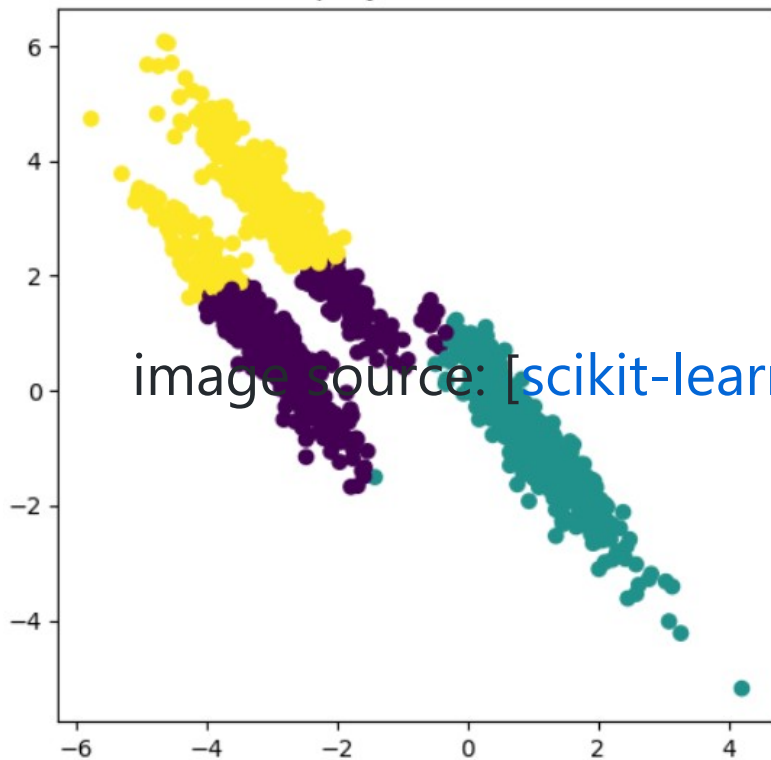
What is  $J'(\text{beer}, \text{bear})$ ?

break time


## discussion

- in the next slide, you are presented with 3 different situations where k-means **didn't work as intended**
- look at the scatter plot and do your best to **explain** why k-means didn't work as intended in each situation
- propose an **approach** for what to do to **avoid** getting in such a trap
- even though the examples are 2 dimensional, your approach should work even when we have more than 2 features and **cannot** rely on **data visualization**

Anisotropically Distributed Blobs



Incorrect Number of Blobs



Unequal Variance





image source: [[scikit-learn.org](https://scikit-learn.org)]



## k-means fail # 1

- we are too **conservative** in our choice of  $k$
- we can catch this by **increasing**  $k$  and noticing a big drop in within-cluster **variability**


image source: [[scikit-learn.org](https://scikit-learn.org)]



## k-means fail # 2

- data distributions follow slanted shapes
- avoid this by excluding **highly correlated** features
- we can try certain **transformations**, e.g. rotation or PCA


image source: [[scikit-learn.org](https://scikit-learn.org)]



## k-means fail # 3

- the middle cluster looks like it should own more of the points around it
- this is a **tough** one
- reruning with different seeds can help identify **borderline points**

image source: [[scikit-learn.org](https://scikit-learn.org)]



**the end**