

# DataSci 520

## lesson 4

### sampling methods



PROFESSIONAL &  
CONTINUING EDUCATION  

---

UNIVERSITY *of* WASHINGTON

## today's agenda

- PDF, CDF and inverse PDF
- common discrete and continuous distributions
- drawing samples from a given distribution
- Monte Carlo methods for estimation

## common discrete distributions

- **discrete uniform:** for equally likely outcomes, such as the result of rolling a dice
- **Bernoulli:** for a single binary outcome, such as a single coin flip
- **binomial:** the number of "successes" in  $n$  independent Bernoulli trials with fixed probability  $p$  of success, where "success" is defined by you, such as the number of heads in  $n = 20$  coin flips
- **poisson:** used for modeling counts, such as the number of customers visiting a store on any day
- **geometric:** the number of Bernoulli trials (with  $p$  fixed) before we see a successful outcome, such as the number of coin flips before we see a tails


## common continuous distributions

- **uniform**: when ranges of data with equal length are equally likely, such as how gas molecules spread in a room
- **normal**: data that is symmetric and bell-shaped, such as people's height, or measurement error if instrument is not **biased**
- **exponential**: heavily right-skewed data, such as lifetime of a light bulb
- **log normal**: skewed data, such as "dwell time" on an online article
- **power law**: for data that appear to follow the "**80-20 rule**"
- **chi-square** is used for a sum of the squares of  $k$  independent standard normal random variables, and is used by many **statistical tests**

let  $A^*$  be the area to the left of  $x^*$

$$A^* = F_X(x^*) = P(X < x^*)$$

which is the CDF function




$x^*$  is drawn from  $X$

$$x^* = F_X^{-1}(A^*)$$

which is the PPF function


# PDF of normal distribution

- the PDF of the **normal distribution** is shown with different means  $\mu$  and standard deviations  $\sigma$
- image source: [wikipedia.org](https://en.wikipedia.org/wiki/Normal_distribution)




# CDF of normal distribution

- the CDF of the **normal distribution** is shown with different means  $\mu$  and standard deviations  $\sigma$
- image source: [wikipedia.org](https://en.wikipedia.org/wiki/Normal_distribution)



## PDF of beta distribution


- the PDF of the **beta distribution** is shown with different parameters  $\alpha$  and  $\beta$
- image source: [wikipedia.org](https://en.wikipedia.org/wiki/Beta_distribution)






## CDF of beta distribution

- the CDF of the **beta distribution** is shown with different parameters  $\alpha$  and  $\beta$
- image source: [wikipedia.org](https://en.wikipedia.org/wiki/Beta_distribution)



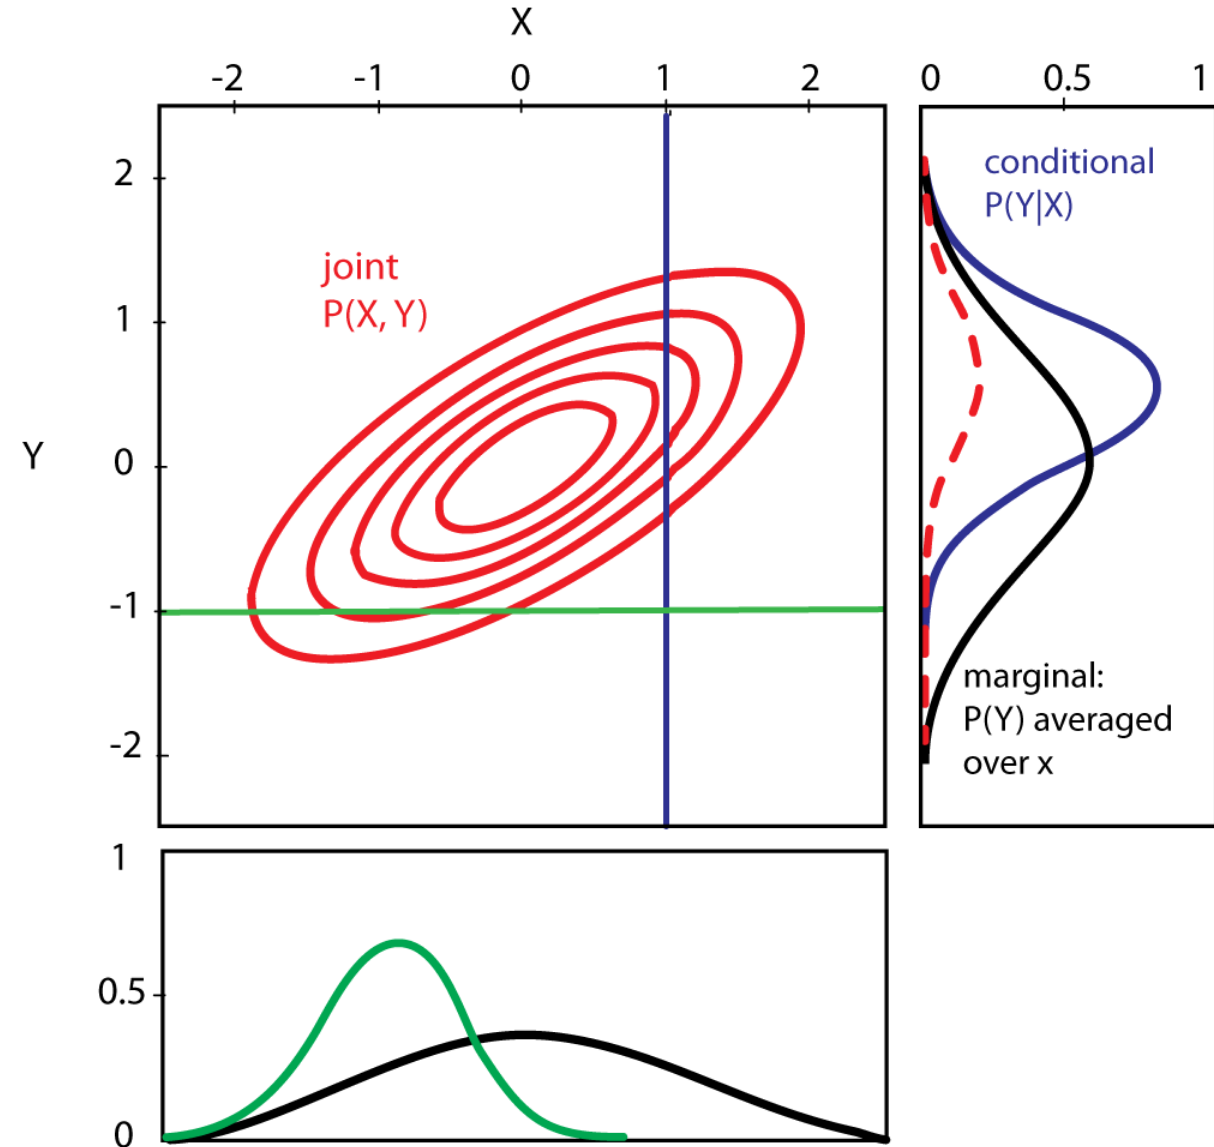
# multivariate normal distribution

- the plot shows the probability density function of  $X$  and  $Y$  together (joint distribution)
- the sample space is 2D:  $(X, Y)$
- the "volume" under the curve adds up to 1
- we can extend this to higher dimensions:  $(X, Y, Z)$  etc.
- image source: [wikipedia.org](https://en.wikipedia.org/wiki/Multivariate_normal_distribution)



# joint density for bivariate normal distribution

- for the multivariate normal distribution, it can be shown that marginal and conditional probabilities also follow a normal distribution
- in general, that it not the case, so the normal distribution is special
- joint and marginals are "fixed", but conditional changes as you slide the blue bar



**notebook time**


**we return to the lecture later**

# drawing from a given distribution

If we know the functional form  $F_X$  of the CDF of  $X$ , we can sample from its distribution using **inverse transform sampling**:

- draw a number from the uniform distribution  $u \sim U(0, 1)$
- we find  $F_X^{-1}(u)$  where  $F_X^{-1}$  is the PPF (which is the inverse of  $F_X$ , the CDF)

This example from [wikipedia.org](https://en.wikipedia.org/wiki/Inverse_transform_sampling) is for the exponential distribution



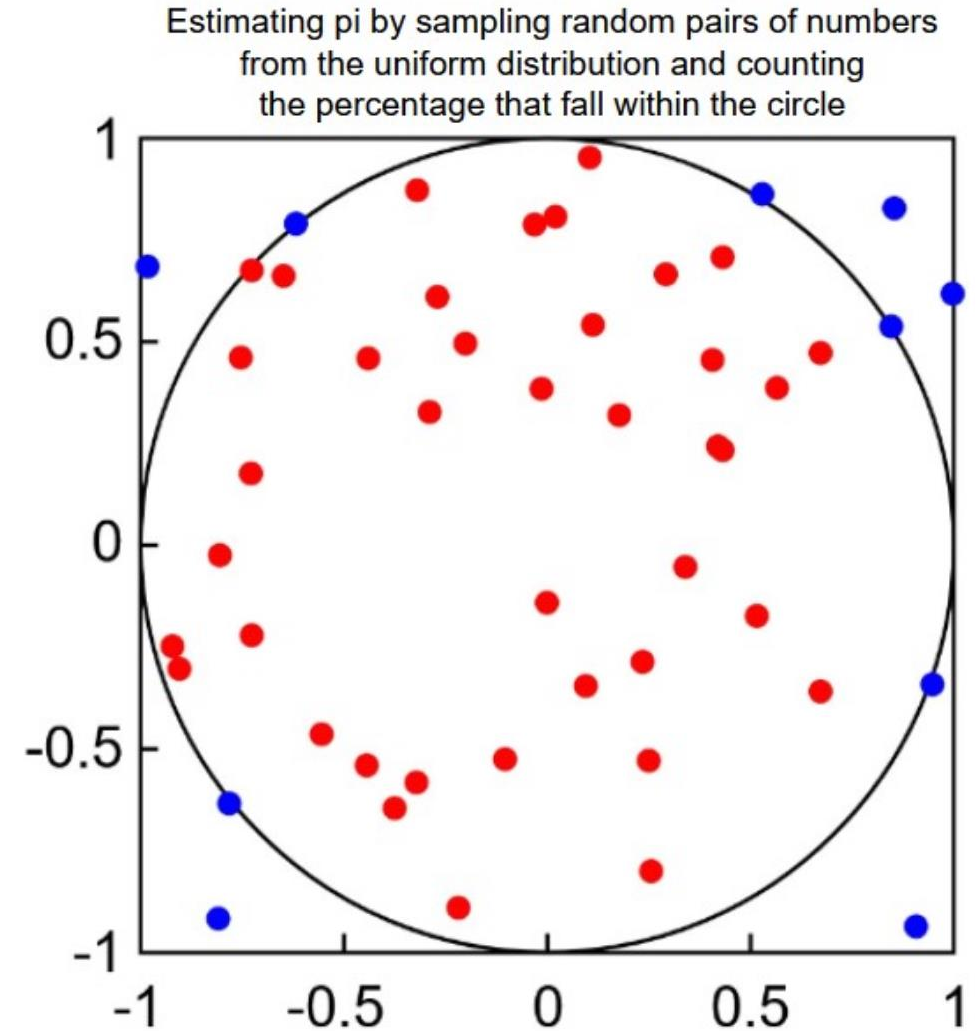
## `numpy` and `scipy` functions

- we can **generate** random numbers from a available distributions
  - in `numpy.random` , for example `np.random.binomial` or `np.random.normal`
  - in `scipy.stats` , for example `scipy.stats.norm.rvs` or `scipy.stats.binom.rvs`
- if we want the CDF, PDF and PPF functions, we go to `scipy.stats` for the distribution and call the corresponding method, for example we have
  - for the geometric distribution `scipy.stats.geom.pmf` , `scipy.stats.geom.ppf` , and `scipy.stats.geom.cdf`
  - for the normal distribution `scipy.stats.norm.pdf` , `scipy.stats.norm.ppf` , and `scipy.stats.norm.cdf`
- Monte Carlo methods is all about generating samples!



# applications of sampling

- **Monte Carlo methods** are used in estimation problems that are hard to solve analytically
- sampling can be used to generate fake data that looks similar to real data (e.g. masking data for privacy reasons)
- sampling can be used by machine learning models to generate new examples, such as in **natural language generation**
- source: [wikipedia.org](https://en.wikipedia.org)




# Monte Carlo Methods

- > Monte Carlo is another name for statistical sampling methods of great importance to physics and computer science
- > Applications of Monte Carlo Method
  - Evaluating integrals of arbitrary functions of 6+ dimensions
  - Predicting future values of stocks
  - Solving partial differential equations
  - Sharpening satellite images
  - Modeling cell populations
  - Finding approximate solutions to NP-hard problems




# Monte Carlo Simulations



# Monte Carlo Simulation Example


> How to evaluate integral of  $f(x)$ ?



$$\int_0^1 f(x) dx = ?$$

# Monte Carlo Simulation Example


- > Can approximate using another function  $g(x)$



$$\int_0^1 f(x) dx = \int_0^1 g(x) dx$$


# Monte Carlo Simulation Example

- > Can approximate by taking the average or expected value



# Monte Carlo Simulation Example

- > Estimate the average by taking N samples



$$\int_0^1 f(x) dx = \frac{1}{N} \sum_{i=1}^N f(x_i)$$

$$E(x) \approx \frac{1}{N} \sum_{i=1}^N x_i$$

# Monte Carlo Integration

$$\mathbf{I} = \int_a^b f(x)dx$$


$$\mathbf{I}_m = (b-a) \frac{1}{N} \sum_{i=1}^N f(x_i)$$

- >  $\mathbf{I}_m$  = Monte Carlo estimate
- >  $N$  = number of samples
- >  $x_1, x_2, \dots, x_N$  are uniformly distributed random numbers between  $a$  and  $b$

$$\lim_{N \rightarrow \infty} \mathbf{I}_m = \mathbf{I}$$

# Monte Carlo Simulation Example

> The variance describes how much the sampled values *vary* from each other.




$$Var[E(f(x))] = \sum_{i=1}^N [f(x_i) - E(f(x))]^2$$

$$Var[E(f(x))] = \frac{1}{N} Var[f(x)]$$

- Variance proportional to  $1/N$

# Monte Carlo Simulation Example

- > Standard Deviation is just the square root of the variance
- > Standard Deviation proportional to  $1 / \sqrt{N}$




- > Need 4X samples to halve the error



# Monte Carlo Simulation Example

> Problem:

- Variance (noise) decreases slowly
- Using more samples only removes a small amount of noise




# Types of Sampling

- > Convenience or Accidental Sampling (This is bad).
  - Grabbing whatever is easier.
  - Grabbing whatever is available/possible.
- > Cluster Sampling
  - Divide the data into clusters and sample select few clusters.
- > Simple Random Sample (Most common)
  - Every point is subjected to a probability of being sampled.
- > Stratified Sampling
  - Sampling subpopulations in a representative fashion.
  - This is sometimes important for class-imbalance problems.
- > Systematic Sampling
  - Sampling every  $k$ -th element of a population. (Common in SQL servers, “table sampling”)
  - Bad for time series data.


# Monte Carlo Simulation Example

- > Importance Sampling: use more samples in important regions of the function
- > If function is high in small areas, use more samples there



# Monte Carlo Simulation Example : Stratified Sampling

- > Partition  $S$  into smaller domains  $S_i$
- > Evaluate integral as sum of integrals over  $S_i$
- > Example: jittering for pixel sampling
- > Often works much better than importance sampling in practice



## Monte Carlo methods example


- recall the example of the random variable  $S$  representing the sum of two independent dice rolls
- we found the distribution of  $S$  by counting all possible outcomes and their relative frequency, for example in 5 out of 36 cases,  $S = 6$ , so  $P(S = 6) = 5/36$
- but what if  $S$  represents the sum of 13 independent dice rolls? there are  $6^{13} = 13,060,694,016$  possibilities, so forget about counting
- instead we **generate samples** from the sampling space, the more samples the better the estimate
- we can then report estimates from the samples or visualize its distribution

```
import numpy as np
import seaborn as sns

n_iter = 10000
s_samples = np.ones(n_iter)

for i in range(n_iter):
    s = 0
    for j in range(13):
        s += np.random.randint(6)
    s_samples[i] = s

sns.distplot(s_samples, bins = 23)
```



## state your assumptions and derivations

- in the previous example, we assume the events (dice rolls) were independent, which simplifies the calculations (and the code) a lot
- in general this is not the case: we can have dependent events so it important to
  - list your constants, random variables, and what you want to estimate
  - for the random variables, state what they represent, what other variables they depend on (are derived from), and what conditional distribution they are assumed to follow given those other variables
  - repeatedly sample from the space in an order that respects the dependencies in the derivations
- for an example, refer to this week's assignment

# Parallelism in Monte Carlo Methods

- > Monte Carlo methods often amenable to parallelism
  - > Find an estimate about  $p$  times faster
- OR
- > Reduce error of estimate by  $p^{1/2}$



# Large Samples and the Law of Large Numbers

- > If we roll a die 60 times and then 600 times, which of the dice will more likely have exactly  $1/6^{\text{th}}$  of the rolls equal to 6 appearing?
  - $P(x=10 | 60 \text{ trials})=?$ 
    - > 0.13701
  - $P(x=100 | 600 \text{ trials})=?$ 
    - > 0.04366
- > Visit: <https://stattrek.com/online-calculator/binomial> for probability calculations.

# Large Samples and the Law of Large Numbers

- > If we roll a die 60 times and then 600 times, which of the dice will more likely have exactly  $1/6^{\text{th}}$  of the rolls equal to 6 appearing?
  - $P(x=10 | 60 \text{ trials})=?$ 
    - >  $\text{BinomPDF}(10) = 0.13701$
  - $P(x=100 | 600 \text{ trials})=?$ 
    - >  $\text{BinomPDF}(100)=0.04366$
- > Which die will be more likely to be within 5% of a  $1/6$ ?
  - $P(7 < x < 13 | 60 \text{ trials})=?$ 
    - >  $\text{BinomCDF}(x=13) - \text{BinomCDF}(x=7)$
    - >  $\text{BinomPDF}(x \leq 13) - \text{BinomPDF}(x \leq 7)$
    - >  $= 0.88478 - 0.19580 = 0.68898$
  - $P(70 < x < 130 | 600 \text{ trials})=?$ 
    - >  $\text{BinomCDF}(x=130) - \text{BinomCDF}(x=70)$
    - >  $\text{BinomPDF}(x \leq 130) - \text{BinomPDF}(x \leq 70)$
    - >  $0.99939 - 0.00038 = 0.99901$

# Law of Large Numbers

- > Sample statistics converge to the population statistics as more unbiased experiments are performed.
  - Example: the mean of 50 coin flips  $(0,1)=(T,H)$  is usually farther away from the true mean of 0.5 than if we did 5,000 coin flips.

# Standard Deviation vs. Standard Error

- > Standard Deviation: A Computed Measure of variability in a sample or population.
  - “My sample values have a standard deviation of XYZ.”
- > Standard Error: Measure of variability in the *statistics* of the sample.
  - “I’ve repeated my experiment many times, and the standard deviation of each of the above standard deviations is small”.
- > For example:
  - Standard deviation: On a sample.
  - Standard Error: Standard deviation of *a set of means* calculated from multiple samples.
    - > You can imagine that the larger my sample, the more confident we can be about the mean.
  - Standard error of a statistic decreases by a rate of  $1/\sqrt{n}$  where  $n$  is your sample size.

**the end**